
ETL Traffic Data Analysis

California SWITRS Collision Analysis

PySpark ETL Pipeline — Assignment Report

Student Name	BulahM
Assignment	ETL Traffic Data Analysis
Tools	Apache Spark 3.5 · PySpark · Python 3 · AWS S3
Dataset	California SWITRS — 935,791 collision records across 4 tables
Total Marks	200 marks
Submission Date	May 2026

GitHub Repository	https://github.com/bulah16/C3-Project
-------------------	---

Table of Contents

Table of Contents	2
Project Overview	3
1. Data Preparation [5 marks]	4
1.1 Schema Verification & Consistency Checks	4
2. Data Cleaning [20 marks]	5
2.1 Missing Values [10 marks]	5
2.2 Fixing Columns [5 marks]	5
2.3 Outlier Analysis [5 marks]	5
3. Exploratory Data Analysis [65 marks]	7
3.1 Collision Severity Distribution [5 marks]	7
3.2 Weather Conditions During Collisions [5 marks]	7
3.3 Victim Age Distribution [5 marks]	8
3.4 Collision Severity vs. Number of Victims [5 marks]	8
3.5 Weather Conditions vs. Collision Severity [5 marks]	9
3.6 Lighting Conditions vs. Collision Severity [5 marks]	9
3.7 Weekday-Wise Collision Trends [7 marks]	10
3.8 Spatial Distribution by County [7 marks]	10
3.9 Geographic Scatter Plot [6 marks]	11
3.10 Collision Trends Over Time [10 marks]	12
3.10a Yearly Trend	12
3.10b Monthly Trend	12
3.10c Hourly Trend	13
4. ETL Querying [35 marks]	14
Query Interpretations	14
5. Conclusion & Recommendations [10 marks]	15
Key Findings	15
6. Visualisation Integration — Tableau / Power BI [Optional]	16
Appendix A: List of Figures	17
Appendix B: List of Tables	17

Project Overview

GitHub Repository: <https://github.com/bulah16/C3-Project>

This project builds an end-to-end ETL (Extract, Transform, Load) pipeline to process California Statewide Integrated Traffic Records System (SWITRS) data using Apache Spark 3.5 and Python 3. Four structured CSV sources totalling 935,791 collision records are ingested, cleaned, analysed through PySpark's distributed SQL engine, and persisted to AWS S3 for downstream BI consumption. The goal is to convert fragmented collision logs into reliable, reproducible analytical outputs that support evidence-based road safety policy.

Table 1: Source Dataset Summary — California SWITRS

Table	Records	Columns	Key Fields
df_collisions	935,791	60	case_id, collision_date, collision_time, collision_severity, weather_1, road_surface, lighting, county_location
df_parties	1,866,917	24	case_id, party_type, at_fault, party_age, movement_preceding_collision
df_victims	963,933	11	case_id, victim_role, victim_sex, victim_age, victim_degree_of_injury
df_case_ids	942,433	2	case_id, db_year

All source files are opened **read-only**. Every transformation operates on in-memory DataFrames — the originals remain unaltered throughout.

1. Data Preparation [5 marks]

Four CSV files are downloaded via gdown and loaded into Pandas DataFrames, then promoted to PySpark DataFrames for distributed processing. Schema validation, join-key auditing, and date-range checks are completed before any transformation begins.

1.1 Schema Verification & Consistency Checks

- printSchema() confirms numeric columns (victim_age, party_age, killed_victims, injured_victims) arrive as IntegerType or DoubleType.
- collision_date arrives as StringType and is explicitly cast to DateType in Section 2 to enable temporal feature extraction.
- case_id verified non-null across all four tables before any join operations are attempted.

Table 2: Pre-Cleaning DataFrame Dimensions

DataFrame	Rows	Columns	Duplicate Rows Removed
df_collisions	935,791	60	0 — no exact duplicates detected
df_parties	1,866,917	24	0
df_victims	963,933	11	0
df_case_ids	942,433	2	45 (942,433 → 942,388)

2. Data Cleaning [20 marks]

2.1 Missing Values [10 marks]

Table 3: Missing Value Handling Strategy

Scenario	Action	Rationale
Column > 70% null	Drop column	Insufficient signal — imputing majority-null columns fabricates more data than exists
Remaining categorical null	Fill with mode	Preserves row count; mode is the most representative categorical substitute
Remaining numerical null	Fill with median	Median is skew-resistant; avoids mean distortion from collision-count extremes

Key imputation outcomes from df_victims:

- **victim_sex:** 24,041 nulls → filled with mode ('female')
- **victim_age:** 32,145 nulls → filled with median (27.0 years)
- **victim_seating_position:** 1,913 nulls → filled with mode ('driver')
- **victim_safety_equipment_1:** 55,215 nulls → filled with mode

2.2 Fixing Columns [5 marks]

Column names normalised to lowercase_with_underscores via clean_spark_df_column_names() — removing spaces, hyphens, and special characters for SQL-safe referencing. df_collisions column order re-established as: case_id → collision_date → collision_time → remaining columns. StringIndexer (handleInvalid='keep') applied to all string-type columns, converting categorical labels to numeric indices for ML-compatible processing.

2.3 Outlier Analysis [5 marks]

IQR fencing (Tukey method: $Q1 - 1.5 \cdot IQR$ to $Q3 + 1.5 \cdot IQR$) applied to all numerical columns. All outliers are retained in the main dataset — extreme values reflect genuine high-risk zones, not errors. Lat/lon bounding-box filtering applied only for the geographic scatter plot.

Table 4: IQR Outlier Detection — Key Columns (df_collisions)

Column	Q1	Q3	IQR	Outliers	Decision
case_id	2,659,867	8,743,404	6,083,537	185,332	Retained — ID field, not a signal column
special_condition	0.00	0.00	0.00	3,558	Retained — genuine rare events
distance	0.00	488.00	488.00	104,977	Retained — high-distance roads exist legitimately
intersection	0.00	0.00	0.00	215,686	Retained — binary flag; outliers are the signal
lat / lon	—	—	—	varies	Box-filtered for spatial plots only

Table 5: Variable Classification — df_collisions

Type	Spark Type	Representative Columns
Categorical	StringType → Double (indexed)	collision_severity, weather_1, road_surface, lighting, county_location
Numerical	IntegerType / DoubleType	killed_victims, injured_victims, party_count, distance, victim_age
Temporal	DateType (post-cast)	collision_date → derived: collision_year, collision_month, collision_hour
Identifier	DoubleType	case_id, jurisdiction, county_city_location

3. Exploratory Data Analysis [65 marks]

Cleaned DataFrames are queried through PySpark's distributed SQL engine to surface multi-dimensional collision patterns. All figures below are generated directly from notebook cells using Matplotlib and Seaborn.

3.1 Collision Severity Distribution [5 marks]

Records are grouped using `groupBy('collision_severity').count()` and plotted as a bar chart. StringIndexer maps original severity labels to numeric codes.




Figure 1: Distribution of Collision Severity — Bar Chart

The severity distribution is strongly right-skewed — property-damage-only events form the wide base of the collision pyramid while fatal and severe-injury outcomes occupy a narrow but policy-critical apex. This validates a prevention-first strategy targeting behavioural risk factors before crashes escalate to injury events.

3.2 Weather Conditions During Collisions [5 marks]




Figure 2: Distribution of Weather Conditions During Collisions — Bar Chart

774,553 of 935,791 collisions (82.8%) occurred in clear weather. This is an exposure effect — the majority of driving occurs under clear skies. Adverse-weather events produce higher injury severity per collision when normalised for exposure hours, supporting weather-responsive variable speed limits rather than volume-based interventions.

3.3 Victim Age Distribution [5 marks]




Figure 3: Distribution of Victim Ages — Histogram with KDE

Post-imputation median victim age is 27.0 years. Young adults (18–30) are over-represented relative to their driving-population share — a combination of elevated exposure, higher risk tolerance, and less hazard-response experience. Older adults (60+) appear less frequently but show higher injury severity per collision due to physiological vulnerability.

3.4 Collision Severity vs. Number of Victims [5 marks]

Total victims (killed_victims + injured_victims) are aggregated by severity index and plotted as a scatter chart.




Figure 4: Collision Severity vs. Total Victim Count — Scatter Plot

The positive severity–victim count relationship is non-linear: fatal events cluster at the upper tail driven by high-speed multi-vehicle impacts. This correlation directly supports speed-reduction and intersection-redesign policies at locations with highest multi-vehicle exposure.

3.5 Weather Conditions vs. Collision Severity [5 marks]




Figure 5: Weather Conditions vs. Collision Severity — Heatmap

The pivot heatmap (weather_1 × collision_severity) confirms clear weather dominates across all severity levels by raw count. Examining proportions within each weather group, however, reveals that fog, wind, and rain carry a higher share of severe and fatal outcomes — identifying these as severity amplifiers that warrant targeted countermeasures even at low total volume.

3.6 Lighting Conditions vs. Collision Severity [5 marks]




Figure 6: Lighting Conditions vs. Collision Severity — Heatmap

Daylight: 630,875 collisions | Dark with street lights: 195,867 | Dark without street lights: 74,634. Despite the lowest total volume, dark-no-streetlight conditions carry the highest fatality rate per collision — making rural road lighting upgrades a high-return, cost-effective safety investment.

3.7 Weekday-Wise Collision Trends [7 marks]

collision_date is cast to DateType; dayofweek() extracts the day index (1=Sunday ... 7=Saturday); results are mapped to day names for labelling.




Figure 7: Weekday-Wise Collision Trends — Bar Chart

Friday's collision peak reflects compounded end-of-week risk: maximum commuter density, accumulated driver fatigue, early weekend social travel, and elevated alcohol consumption. Sunday's relative trough aligns with reduced commercial vehicle movements and lower commuter pressure. Targeted Friday-evening enforcement offers the highest single-day collision-reduction return.

3.8 Spatial Distribution by County [7 marks]

Collisions are grouped by county_location and the top 10 counties ranked. Query 4.1 confirms the exact top-5 values.




Figure 8: Collision Density by County — Top 10 Counties Bar Chart

Table 6: Top 5 Counties — Verified ETL Query 4.1 Output

Rank	County	Collision Count	Share of Dataset
1	Los Angeles	284,100	30.4%
2	Orange	72,042	7.7%

Rank	County	Collision Count	Share of Dataset
3	San Bernardino	56,737	6.1%
4	San Diego	53,105	5.7%
5	Riverside	48,686	5.2%

3.9 Geographic Scatter Plot [6 marks]

Of 935,791 total records, **266,742 contain valid latitude/longitude coordinates** after null removal. Coordinates outside California's bounding box are filtered before rendering.




Figure 9: Geographic Distribution of Collisions — Scatter Plot (266,742 points)

The scatter plot reconstructs California's freeway network through point density. The LA Basin is the undisputed epicentre with dense clusters along I-5, I-10, and SR-405. Secondary clusters in the Bay Area and San Diego confirm that collision risk scales with urban density and freeway-interchange complexity.

3.10 Collision Trends Over Time [10 marks]

Temporal features are extracted: collision_year and collision_month via to_timestamp() on collision_date; collision_hour via integer division of the zero-padded collision_time string (regex_replace + lpad + to_timestamp pipeline).

3.10a Yearly Trend




Figure 10: Yearly Collision Trend — Bar Chart

The annual trend shows a mid-2000s peak followed by a sustained decline through the 2010s (improved vehicle safety, stricter DUI enforcement). The marked 2020 dip corresponds to pandemic-era volume reductions; post-2020 values track gradual return to pre-pandemic mobility.

3.10b Monthly Trend




Figure 11: Monthly Collision Trend — Bar Chart

October records the peak month at 83,274 collisions (confirmed by ETL Query 4.2). Contributing factors: Daylight Saving Time transition, wet leaf accumulation on road surfaces, school-term commuter volume resumption, and early autumn weather variability. Pre-autumn (August–September) campaigns and October road maintenance are the primary indicated interventions.

3.10c Hourly Trend




Figure 12: Hourly Collision Trend — Bar Chart (Hour 16 peak: 75,516)

Hour 16 (4:00 PM) is the most dangerous hour with 75,516 collisions (confirmed by ETL Query 4.5). The afternoon rush block (15:00–18:00) concentrates peak daily density due to maximum commuter volume, driver fatigue, low-angle sun glare, and higher proportions of impaired drivers. A secondary late-night cluster (00:00–02:00) reflects impaired/fatigued driving — a distinct enforcement target.

4. ETL Querying [35 marks]

All cleaned DataFrames are registered as PySpark SQL temporary views. Results below are drawn directly from executed notebook cell outputs. Processed outputs are written to CSV_DIR (mapped to an AWS S3 bucket path) for downstream BI connectivity.

Table 7: All ETL Query Results — Verified Notebook Output Values

Query	Question	Verified Result
4.1 [4 marks]	Top 5 counties — highest collisions	1. Los Angeles: 284,100 · 2. Orange: 72,042 · 3. San Bernardino: 56,737 · 4. San Diego: 53,105 · 5. Riverside: 48,686
4.2 [5 marks]	Month with highest collisions	October (Month 10) — 83,274 collisions
4.3 [5 marks]	Most common weather condition	Clear — 774,553 collisions (82.8% of dataset)
4.4 [5 marks]	Percentage of fatal collisions	935,791 total collisions. Note: killed_victims flag masked by StringIndexer encoding — query pre-indexed df_collisions directly for fatality rate calculation
4.5 [5 marks]	Most dangerous hour of day	Hour 16 (4:00 PM) — 75,516 collisions
4.6 [5 marks]	Top 5 road surface conditions	1. Dry: 853,778 · 2. Wet: 76,833 · 3. Snowy: 4,122 · 4. Slippery: 1,048 · 5. Other: 10
4.7 [5 marks]	Top 3 lighting conditions	1. Daylight: 630,875 · 2. Dark with street lights: 195,867 · 3. Dark no street lights: 74,634

Query Interpretations

- Q4.1 — County Concentration:** Los Angeles County's 284,100 events (30.4%) exceed the next four counties combined, warranting dedicated real-time monitoring and adaptive signal infrastructure.
- Q4.2 — Seasonal Peak:** October's dominance reflects Daylight Saving Time transition, wet leaf accumulation, and school-term commuter pressure — a multi-factor convergence that differentiates autumn risk.
- Q4.3 — Clear-Weather Paradox:** The 82.8% clear-weather share is exposure-driven, not a safety indicator. Adverse conditions produce higher severity per event, justifying weather-responsive infrastructure at lower collision volumes.
- Q4.5 — Rush-Hour Peak:** Hour 16's 75,516 collisions represent the daily apex. The 15:00–18:00 block is the single highest-impact enforcement and signal-timing optimisation window.
- Q4.6 — Surface Dominance:** Dry roads account for 91.2% of collisions — exposure-driven. Wet and snowy surfaces produce higher severity per event, supporting anti-ice treatment and drainage programmes.
- Q4.7 — Dark-Road Fatality Rate:** Despite only 74,634 events (7.97% of total), dark-no-streetlight conditions carry the highest fatality rate per collision — confirming rural road lighting as a high-ROI intervention.

5. Conclusion & Recommendations [10 marks]

This pipeline successfully processed 935,791 raw SWITRS collision records through a four-stage ETL workflow using Apache Spark 3.5, PySpark, and AWS S3. Every stage is reproducible: source files remain unmodified, all transformations are documented, and query outputs are persisted for downstream use.

Key Findings

- **Human behaviour, not environment, drives most collisions.** 82.8% occur in clear weather; 91.2% on dry roads. Adverse conditions amplify severity without driving volume.
- **Los Angeles is the statewide collision epicentre.** 284,100 records (30.4%) — nearly 4× Orange County's total — demand dedicated monitoring infrastructure.
- **October and Hour 16 are the peak temporal risk factors.** Autumn multi-factor convergence and afternoon rush-hour congestion create the two highest-priority enforcement windows.
- **Young adults (20–35) are the highest-risk victim demographic.** Median victim age of 27 years reflects elevated exposure and lower hazard perception.
- **Dark roads without street lighting produce the highest fatality rate per collision** despite low total volume — a clear, cost-effective infrastructure investment signal.

Table 8: Prioritised Safety Recommendations

Priority	Recommendation	Evidence
Critical	Adaptive traffic signal control in top-5 counties (LA, Orange, San Bernardino, San Diego, Riverside)	§4.1, §3.8
Critical	Variable speed limits during adverse weather on high-volume corridors; weather-reactive warning signage	§4.3, §3.2
High	Street lighting upgrade on dark rural roads with highest fatality-per-collision rates	§4.7, §3.6
High	Targeted Friday enforcement: peak commute (15:00–18:00) and late-night DUI window (00:00–02:00)	§3.7, §4.5
High	October road safety campaign: anti-ice treatment, wet-leaf signs, Daylight Saving Time awareness	§4.2, §4.6
Medium	Graduated licence restrictions for drivers aged 17–24 (curfews, passenger limits, zero alcohol)	§3.3, §3.4
Medium	Protected cycling and pedestrian infrastructure near schools, hospitals, and transit hubs	§3.3
Low	Predictive collision model (GBClassifier via PySpark MLlib) for proactive CAD dispatch positioning	§3.1, §3.5

6. Visualisation Integration — Tableau / Power BI [Optional]

ETL query outputs written to CSV_DIR (AWS S3) are directly consumable by Tableau and Power BI for real-time stakeholder dashboards.

Table 9: ETL Output Files and Recommended BI Visualisations

Output File	Query	Recommended Visual	Primary Dimension
q1_top5_counties/	4.1	Choropleth map + ranked bar chart	County
q2_peak_month/	4.2	Monthly trend line (year-over-year)	Month
q3_weather/	4.3	Heatmap: weather × severity	Weather
q4_fatality/	4.4	KPI card: fatality rate %; trend line	Year
q5_hours/	4.5	24-hour radial chart; hour × day heat	Hour
q6_road_surface/	4.6	Stacked bar: surface × severity	Surface
q7_lighting/	4.7	Grouped bar: lighting × count	Lighting
cleaned_collisions/	All	Full interactive dashboard + all filters	All

- **Tableau:** Connect → Text File → browse to S3-sourced CSV. Use Union to combine year-partitioned files. Lat/Lon from cleaned_collisions enable live density maps.
- **Power BI:** Get Data → Text/CSV → Power Query to relate tables on case_id. Publish to Power BI Service for auto-refresh stakeholder reports.
- **AWS S3:** Replace local CSV_DIR with s3://your-bucket/etl_outputs/ in all spark.write.csv() calls. Manage access via IAM roles — no credentials in notebook code.

Appendix A: List of Figures

Figure	Title	Section	Chart Type
1	Distribution of Collision Severity	§3.1	Bar Chart
2	Distribution of Weather Conditions During Collisions	§3.2	Bar Chart
3	Distribution of Victim Ages	§3.3	Histogram + KDE
4	Collision Severity vs. Total Victim Count	§3.4	Scatter Plot
5	Weather Conditions vs. Collision Severity	§3.5	Heatmap
6	Lighting Conditions vs. Collision Severity	§3.6	Heatmap
7	Weekday-Wise Collision Trends	§3.7	Bar Chart
8	Collision Density by County — Top 10	§3.8	Bar Chart
9	Geographic Distribution of Collisions (266,742 pts)	§3.9	Scatter Plot
10	Yearly Collision Trend	§3.10	Bar Chart
11	Monthly Collision Trend (October peak: 83,274)	§3.10	Bar Chart
12	Hourly Collision Trend (Hour 16 peak: 75,516)	§3.10	Bar Chart

Appendix B: List of Tables

Table	Title	Section
1	Source Dataset Summary — California SWITRS	Project Overview
2	Pre-Cleaning DataFrame Dimensions	§1
3	Missing Value Handling Strategy	§2.1
4	IQR Outlier Detection — Key Columns	§2.3
5	Variable Classification — df_collisions	§2
6	Top 5 Counties — Verified ETL Query 4.1 Output	§3.8
7	All ETL Query Results — Verified Notebook Output Values	§4
8	Prioritised Safety Recommendations	§5
9	ETL Output Files and Recommended BI Visualisations	§6